

SESSION 05

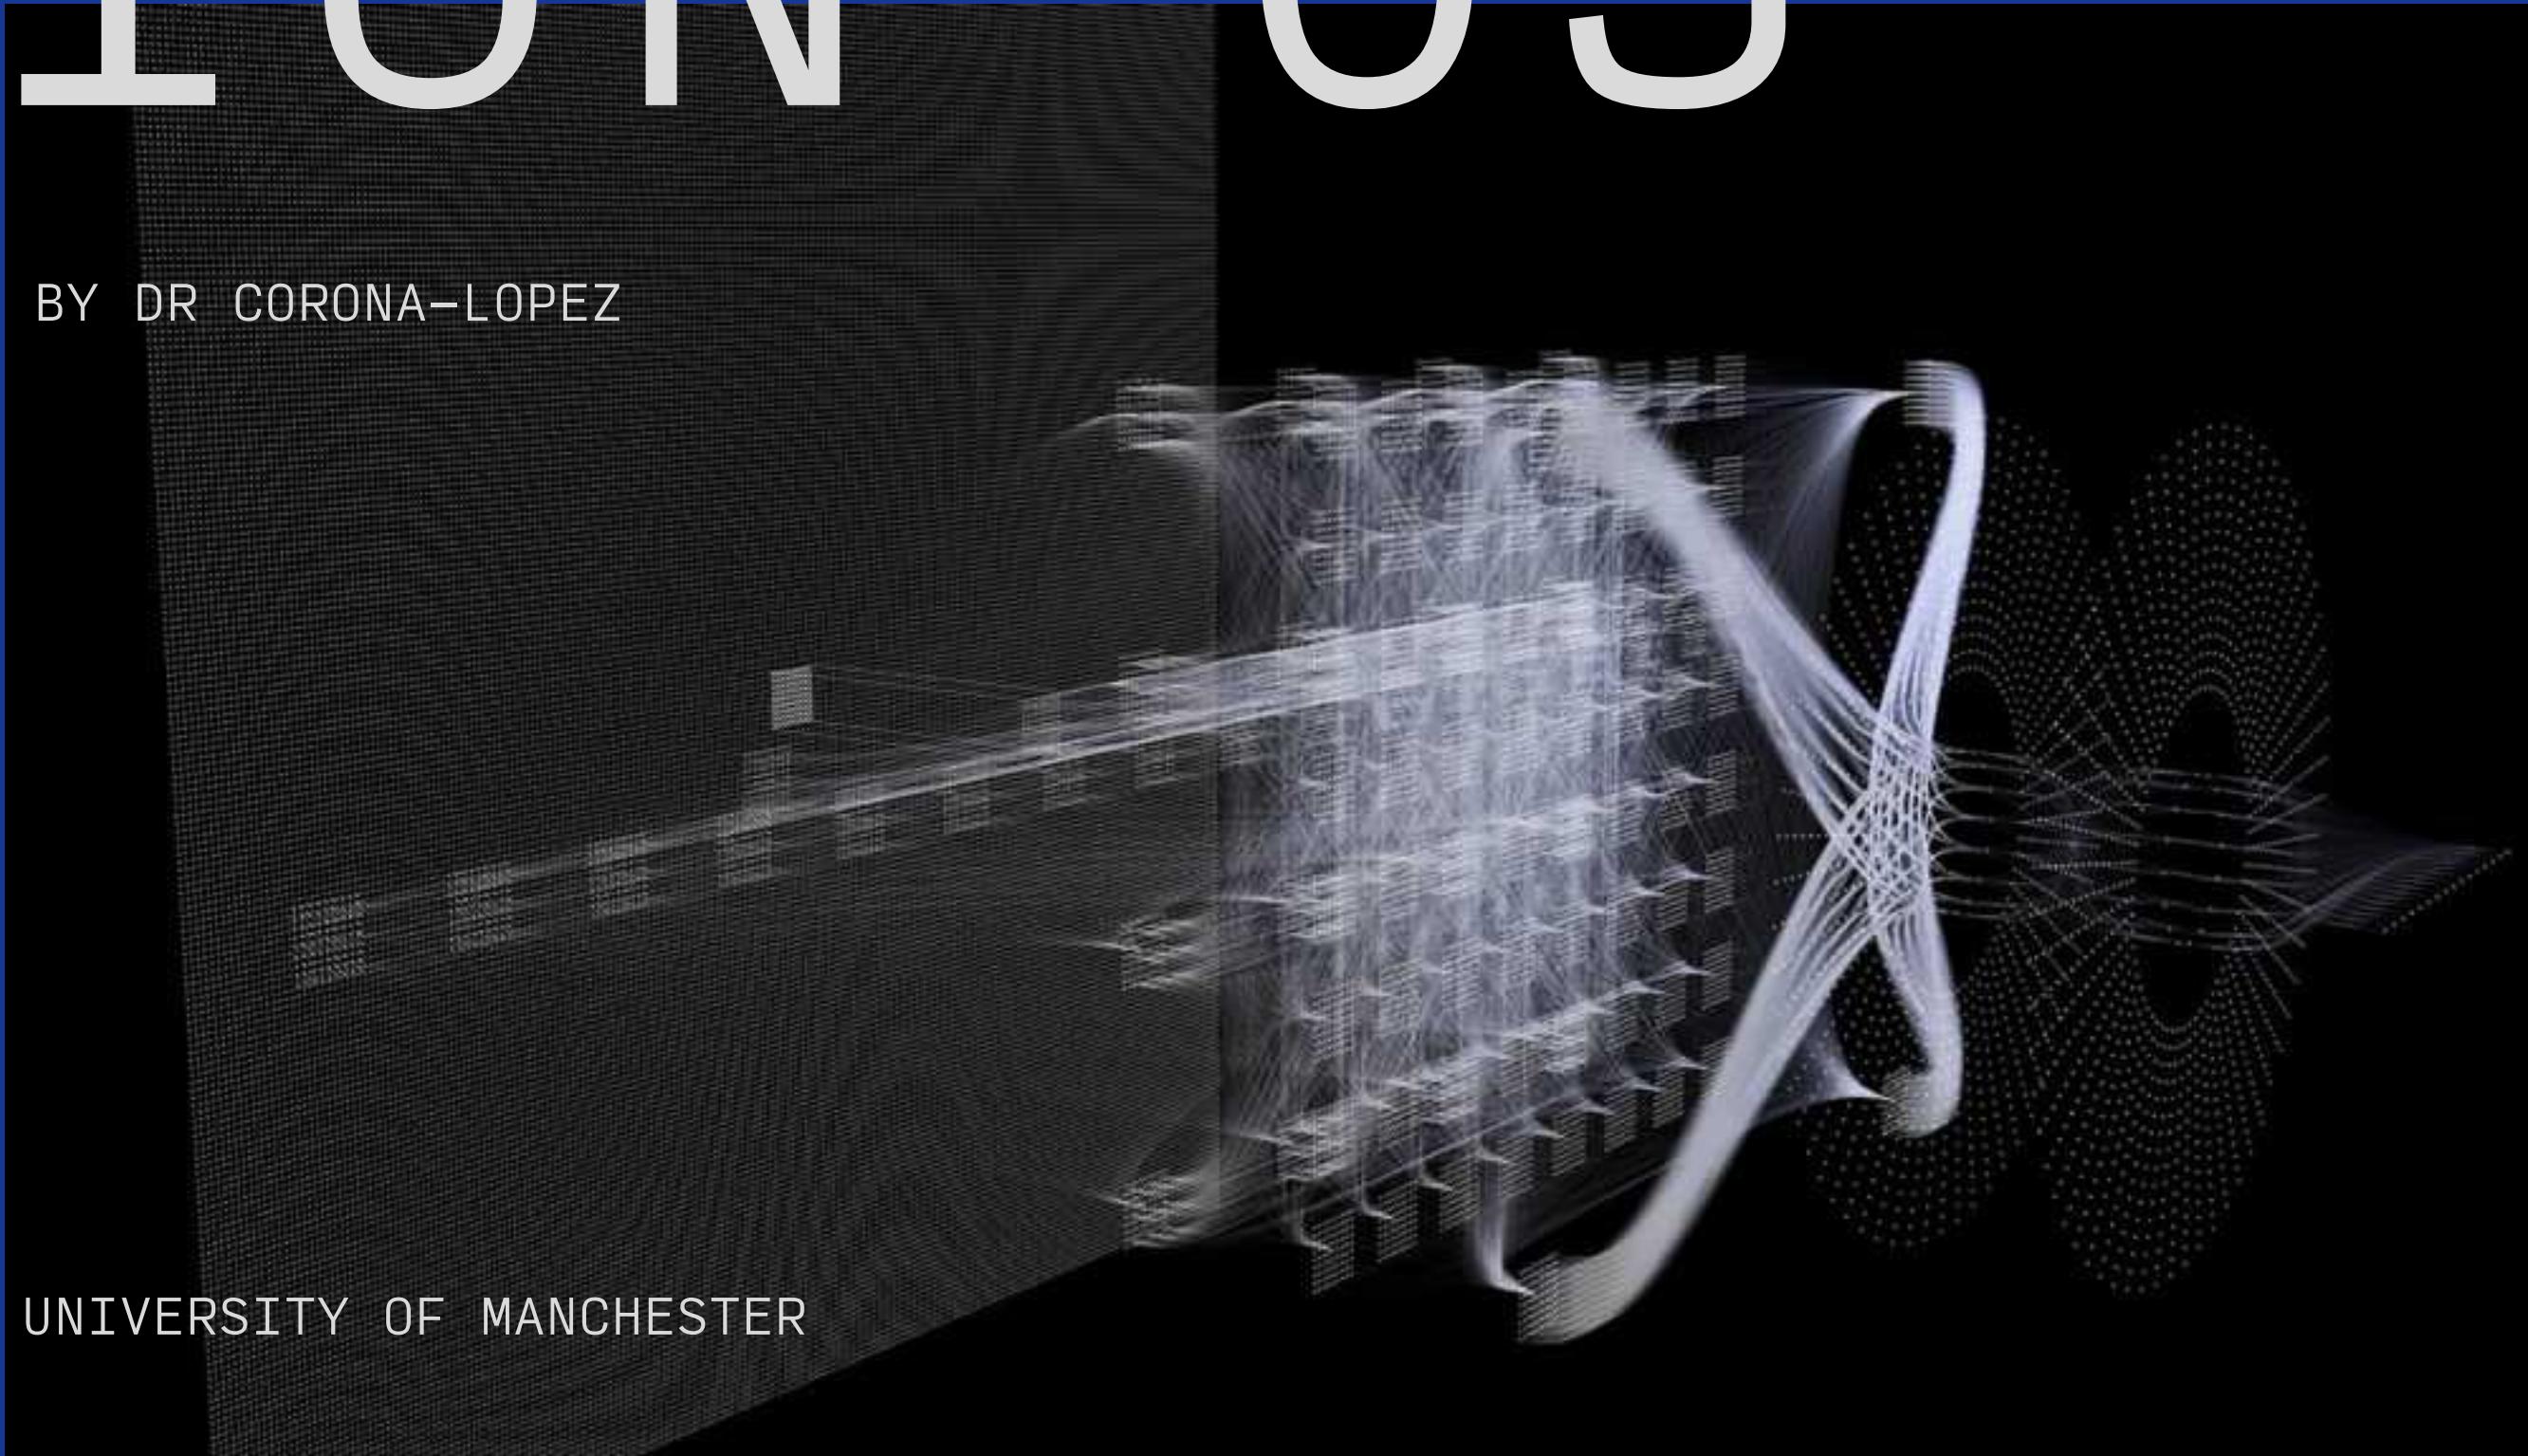
WORKSHOP

MARCH 2026

TRANSFER LEARNING

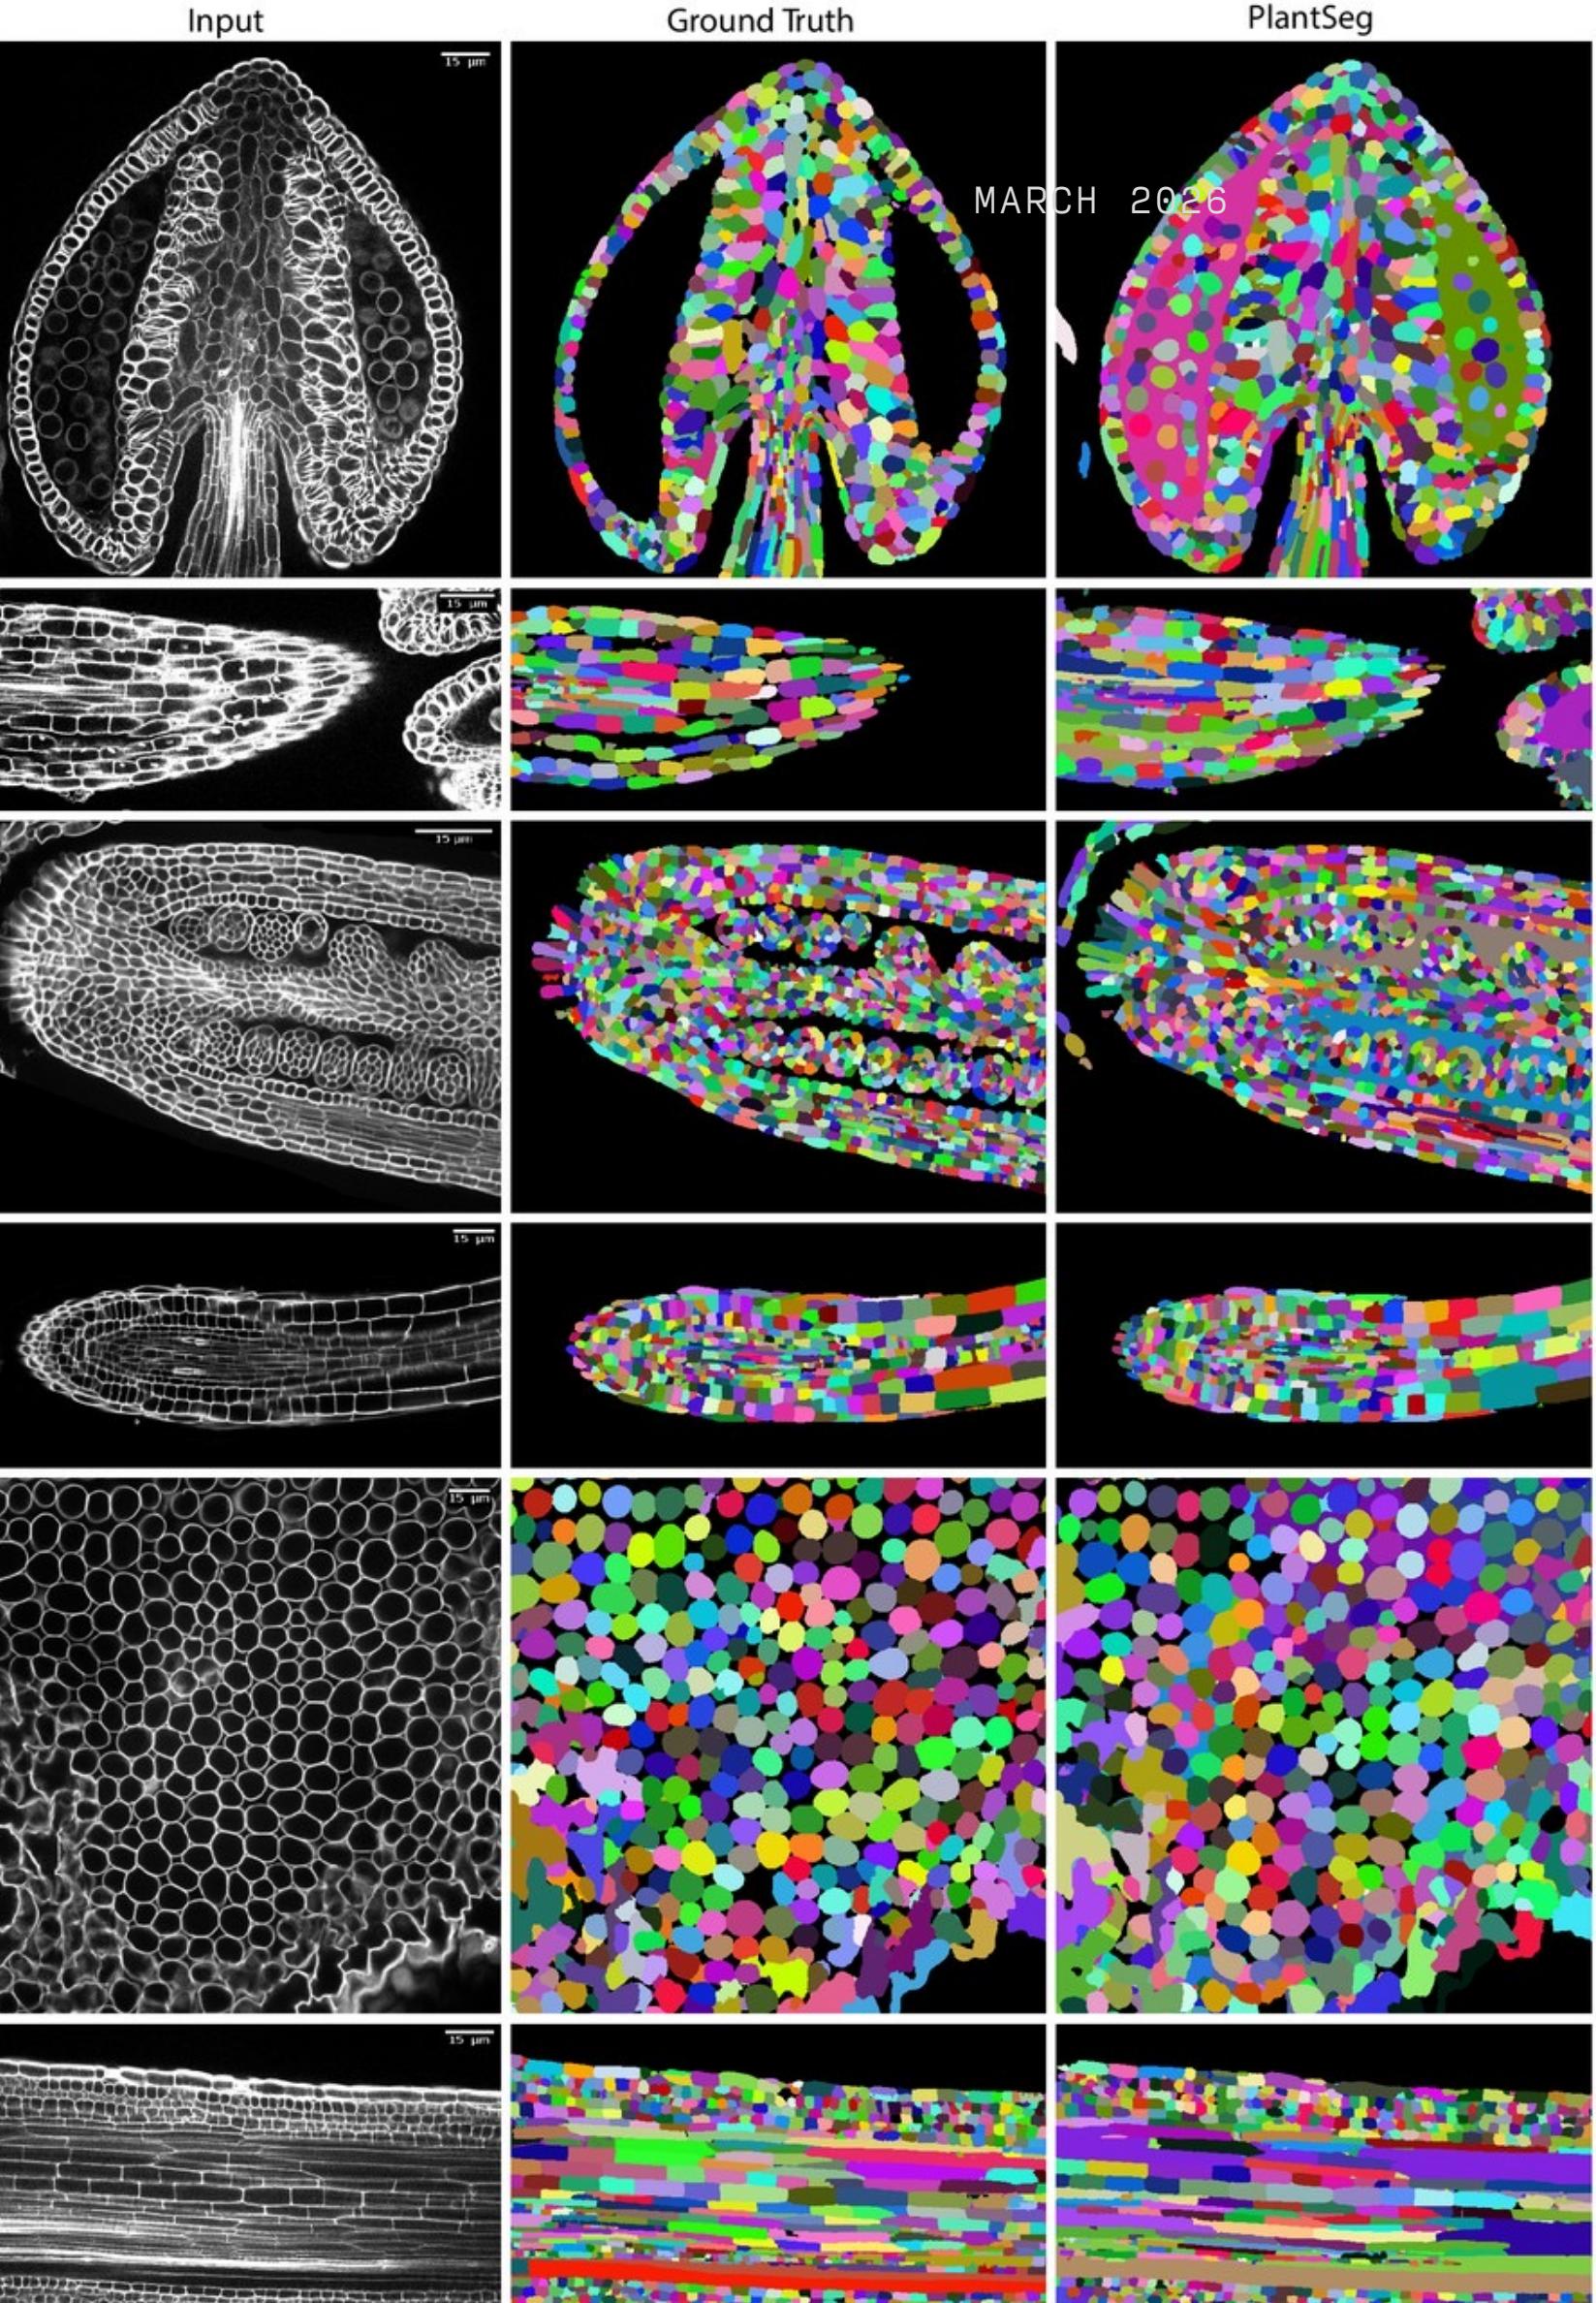
FACULTY OF SCIENCE AND ENGINEERING

UNIVERSITY OF MANCHESTER



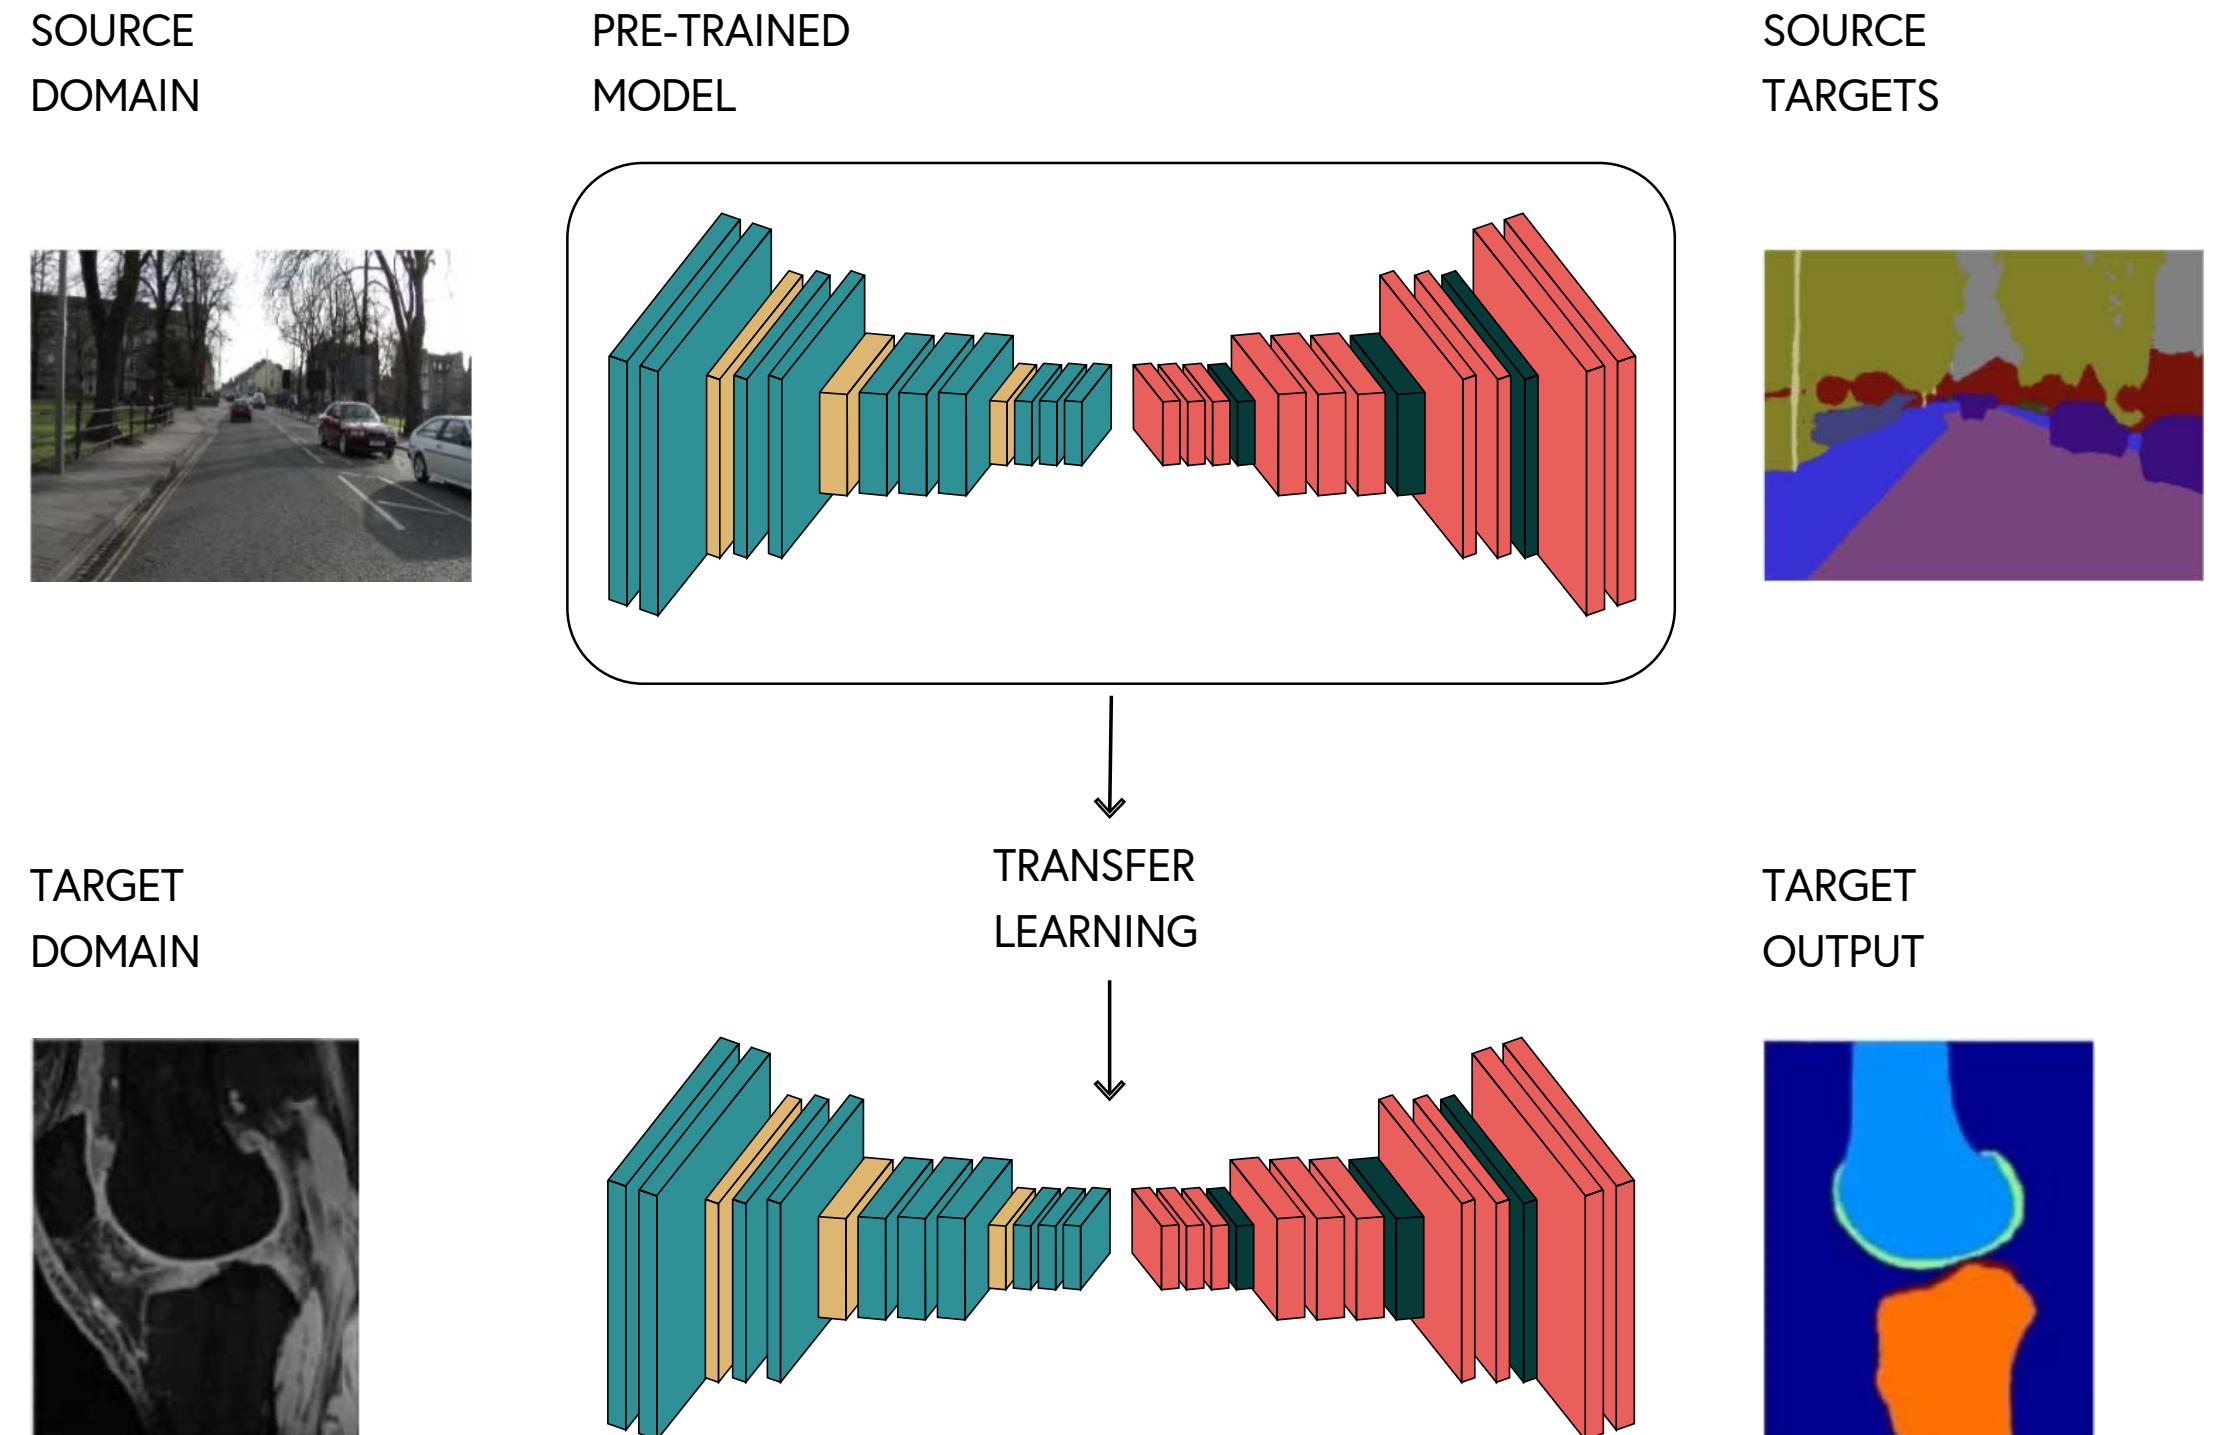
AGENDA

- **transfer learning:** core concepts, and when to use it
- **dataset preparation:** custom PyTorch dataset structure, computing dataset stats for normalization
- **albumentations:** synchronised image & mask augmentation
- **Baseline architecture & loss:** intro to U-Net structure, and tackling imbalance with dice loss
- **implementing transfer learning:** intro to EfficientNet architecture, building Efficient U-Net, handling feature map mismatches, and freezing pre-trained weights



transfer learning is a technique where a model developed for one task is reused as a starting point for a model on a second task.

it leverages knowledge from pre-trained models instead of starting from scratch
it is particularly effective for deep learning models that require massive datasets and computational resources



USE CASES

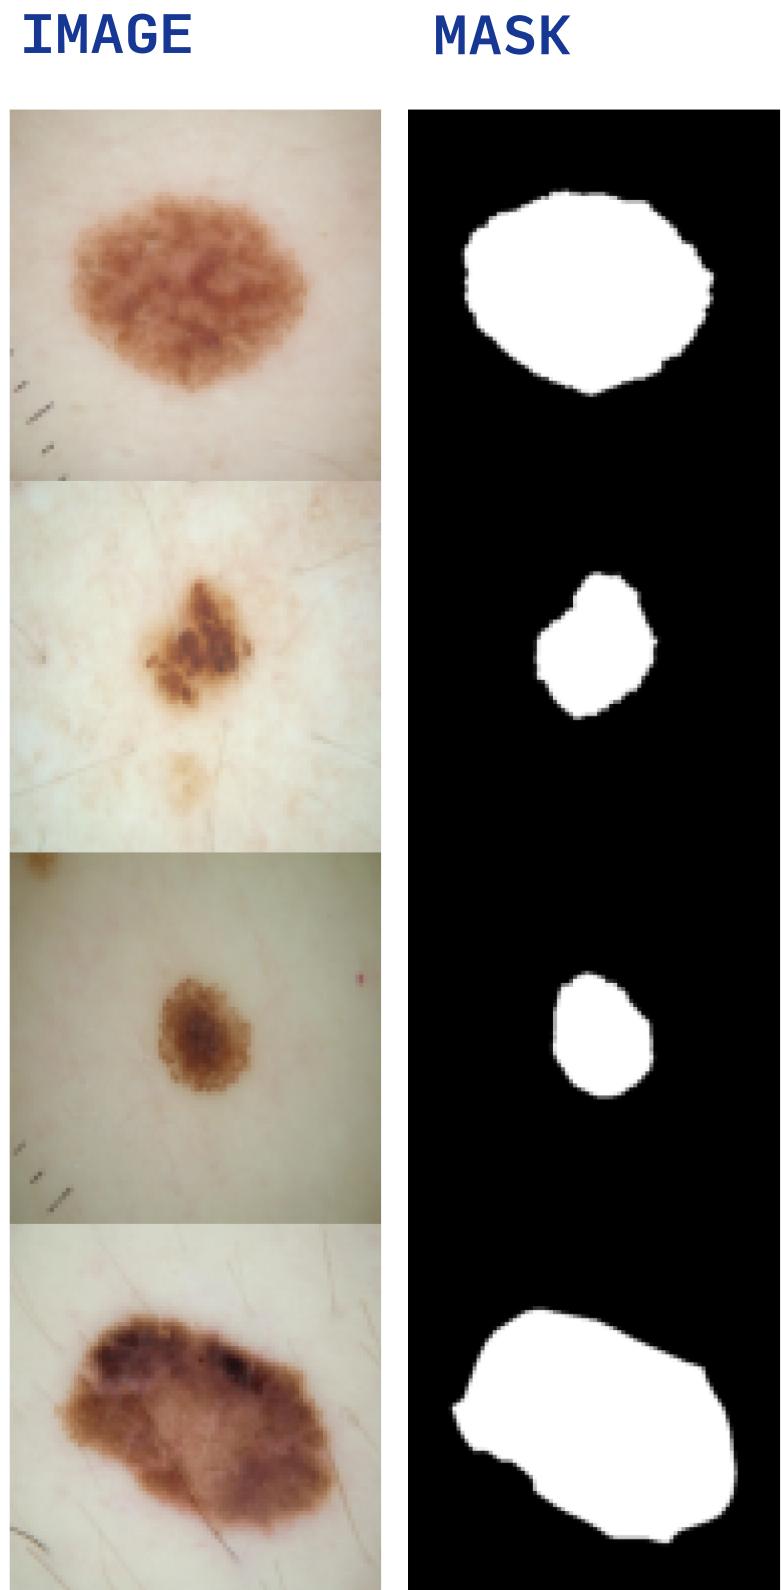
the **effectiveness** depends on similarity between source and target domains

| SCENARIO | EXAMPLE | BENEFIT |
|------------------------|--|---|
| LIMITED TRAINING DATA | medical imaging with few samples | pre-trained features compensate for data scarcity |
| SIMILAR DOMAINS | from natural images to satellite imagery | underlying features (edges, textures) transfer well |
| TIME CONSTRAINTS | rapid prototyping needs | accelerates model development cycle |
| HARDWARE LIMITATIONS | training with limited gpu access | reduces computational requirements |
| PREVENTING OVERFITTING | small dataset applications | regularization effect from pre-trained weights |

ISIC 2016

medical datasets like ISIC are typically smaller than general computer vision datasets, making transfer learning particularly valuable

- contains dermoscopic images of skin lesions with expert-annotated segmentation masks
- 900 training images and 379 test images with corresponding binary masks
- critical for developing automated diagnostic tools for early melanoma detection
- challenging due to varying lesion sizes, shapes, colours, and skin types



DATA NORMALIZATION



```
# Conceptual loop for computing statistics
channel_sum = torch.tensor([0.0, 0.0, 0.0])
channel_squared_sum = torch.tensor([0.0, 0.0, 0.0])
num_batches = 0

for images, _ in dataloader:
    # images shape: [batch_size, 3, height, width]
    channel_sum += torch.mean(images, dim=[0, 2, 3])
    channel_squared_sum += torch.mean(images ** 2, dim=[0, 2, 3])
    num_batches += 1

mean = channel_sum / num_batches
std = (channel_squared_sum / num_batches - mean ** 2) ** 0.5
```

pixel values typically range from 0 to 255. neural networks struggle with large, unscaled inputs, leading to slow and unstable training.

the solution: we normalize the dataset using z-score normalization: $x_{normalized} = \frac{x-\mu}{\sigma}$

while we can use the generic imagenet statistics (mean: [0.485, 0.456, 0.406]), medical images (like pink/brown skin lesions) have a completely different color distribution.

calculating the exact dataset statistics yields better results.

how it works: we iterate over the entire training set to calculate the sum and squared sum of pixel values across the red, green, and blue channels.

DATA AUGMENTATION & SYNC



```
import albumentations as A

# Albumentations handles BOTH image and mask simultaneously
transform = A.Compose([
    A.HorizontalFlip(p=0.5),
    A.VerticalFlip(p=0.5),
    A.RandomRotate90(p=0.5)
])

# Usage in __getitem__:
# augmented = transform(image=image_array, mask=mask_array)
```

in standard image classification, if you rotate a picture of a cat 90 degrees, the label remains "cat".

the segmentation way (sync): in segmentation, the "label" is a pixel-perfect mask. if you flip the image, you must flip the mask in the exact same way. if they fall out of sync, the network learns garbage.

the solution: standard torchvision struggles with synced mask augmentation.

this is why we switch to the albumentations library, which handles synchronized (image, mask) transforms natively.

DATA AUGMENTATION & SPLITTING



```
# Inside the updated ISICDataset __getitem__ method:  
    if self.transform:  
        # Albumentations expects named arguments for sync transforms!  
        augmented = self.transform(image=image_np, mask=mask_np)  
        image = augmented['image']  
        mask = augmented['mask']  
  
# --- Splitting the Data ---  
from torch.utils.data import random_split  
  
total_size = len(full_dataset)  
train_size = int(0.8 * total_size)  
val_size = total_size - train_size  
  
train_dataset, val_dataset = random_split(full_dataset,  
                                         [train_size, val_size])
```

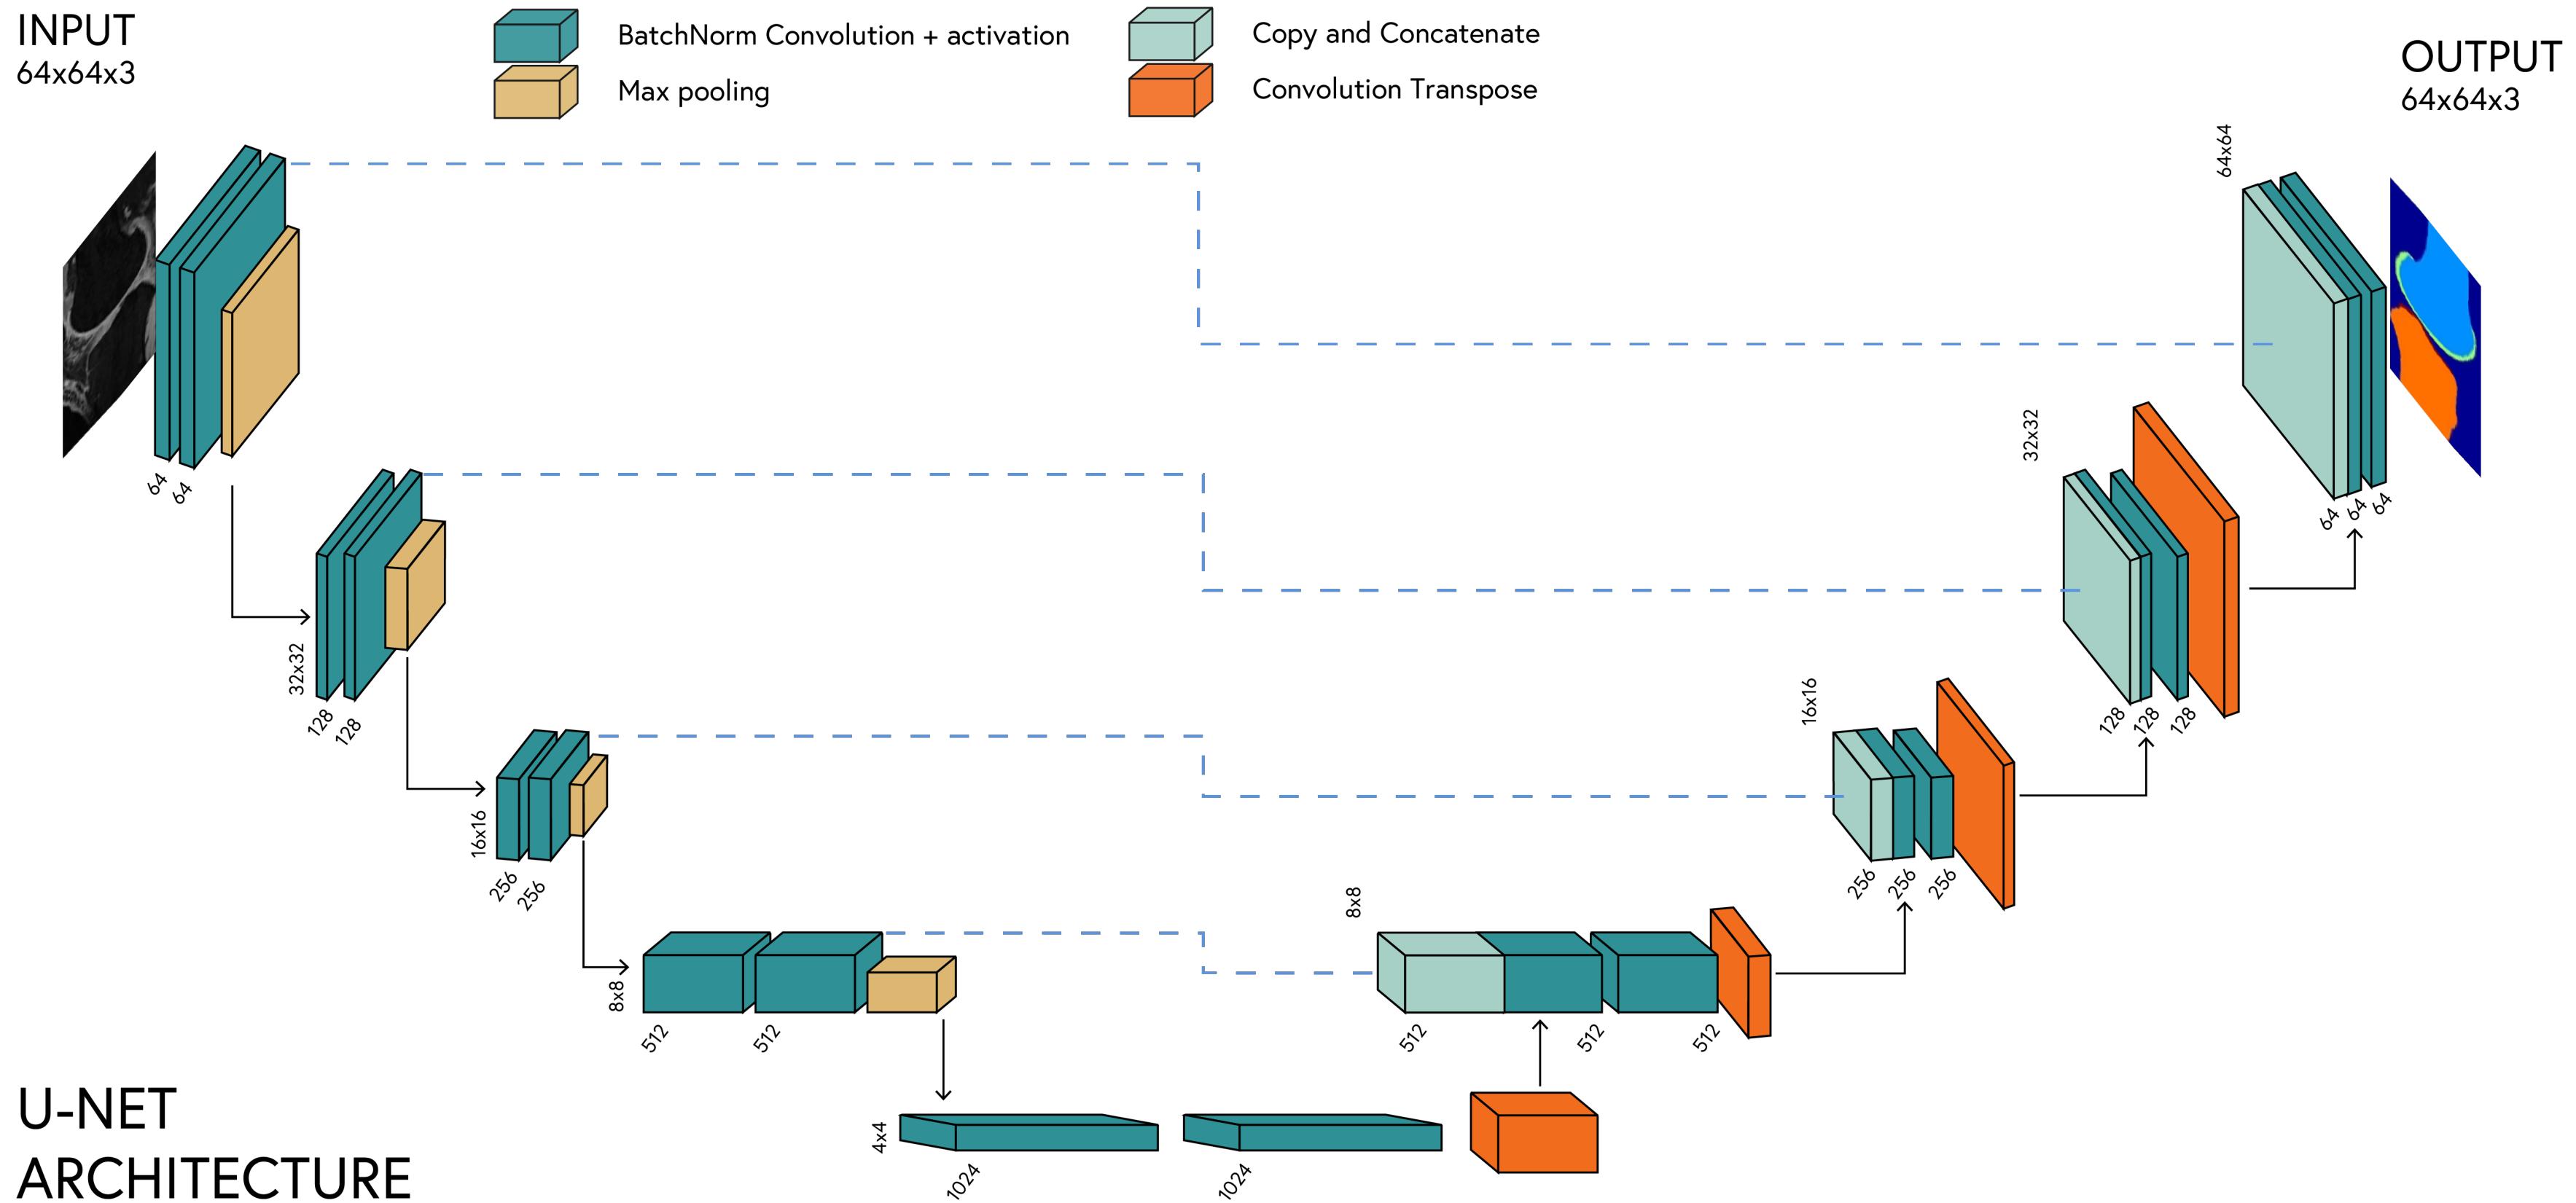
the wrapper: because albumentations requires numpy arrays and returns a dictionary, we need to update our dataset's `__getitem__` method to properly pass both the image and mask to the transform function.

train/validation split: we never train on all our data. we must reserve a portion (usually 20%) to validate that our model is actually learning generalizable features, not just memorizing the training images.

pytorch random_split: a built-in utility that randomly divides the dataset into non-overlapping subsets.

U-NET

- **encoder-decoder architecture** with **skip connections**
- captures both context and localisation information
- widely used for biomedical image segmentation



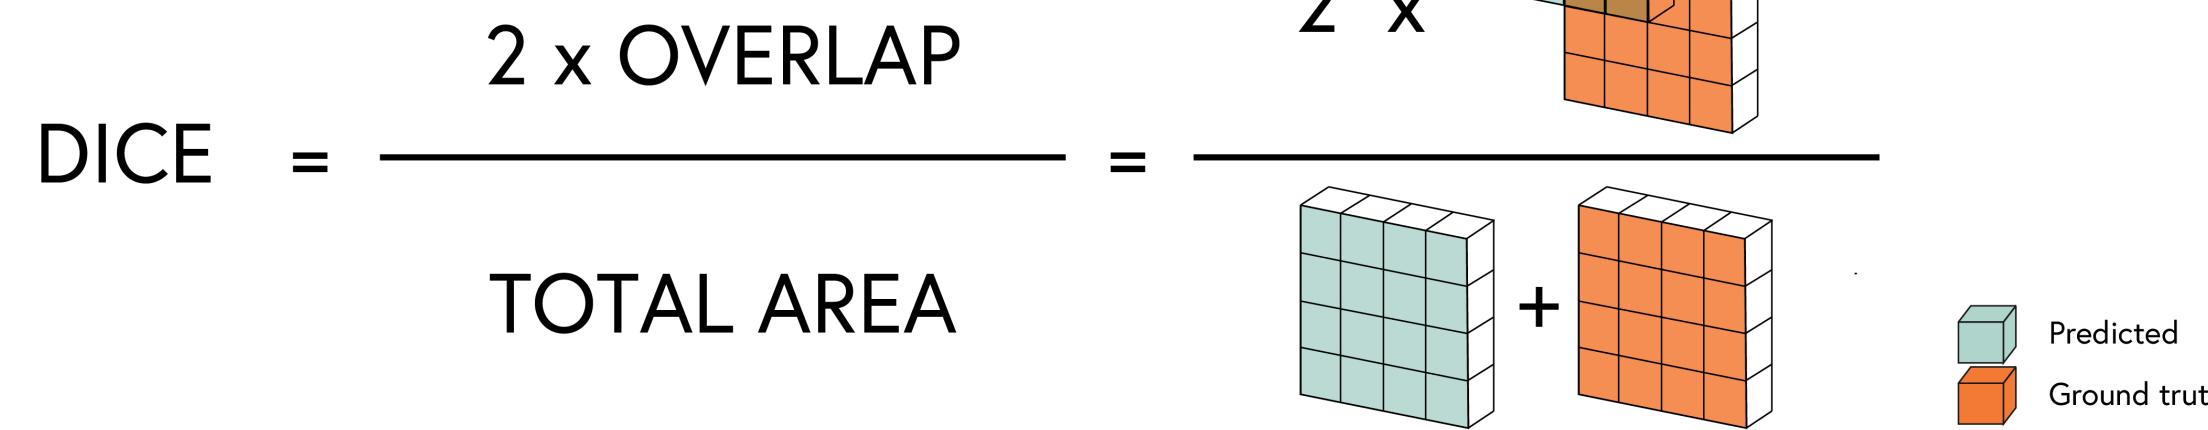
DICE LOSS

the class imbalance problem: in medical imaging, the region of interest (e.g., a skin lesion) often occupies only 5% of the image, while healthy skin/background is 95%.

dice loss:

- measures **overlap** between predicted and ground truth masks
- handles **class imbalance** better than binary cross-entropy

$$\text{DiceLoss} = 1 - \frac{2|X \cap Y|}{|X| + |Y|}$$



```

● ● ●

class DiceLoss(nn.Module):
    def __init__(self, smooth=1e-6):
        super().__init__()
        self.smooth = smooth # Prevents division by zero

    def forward(self, pred, target):
        # Flatten the tensors
        pred = pred.view(-1)
        target = target.view(-1)

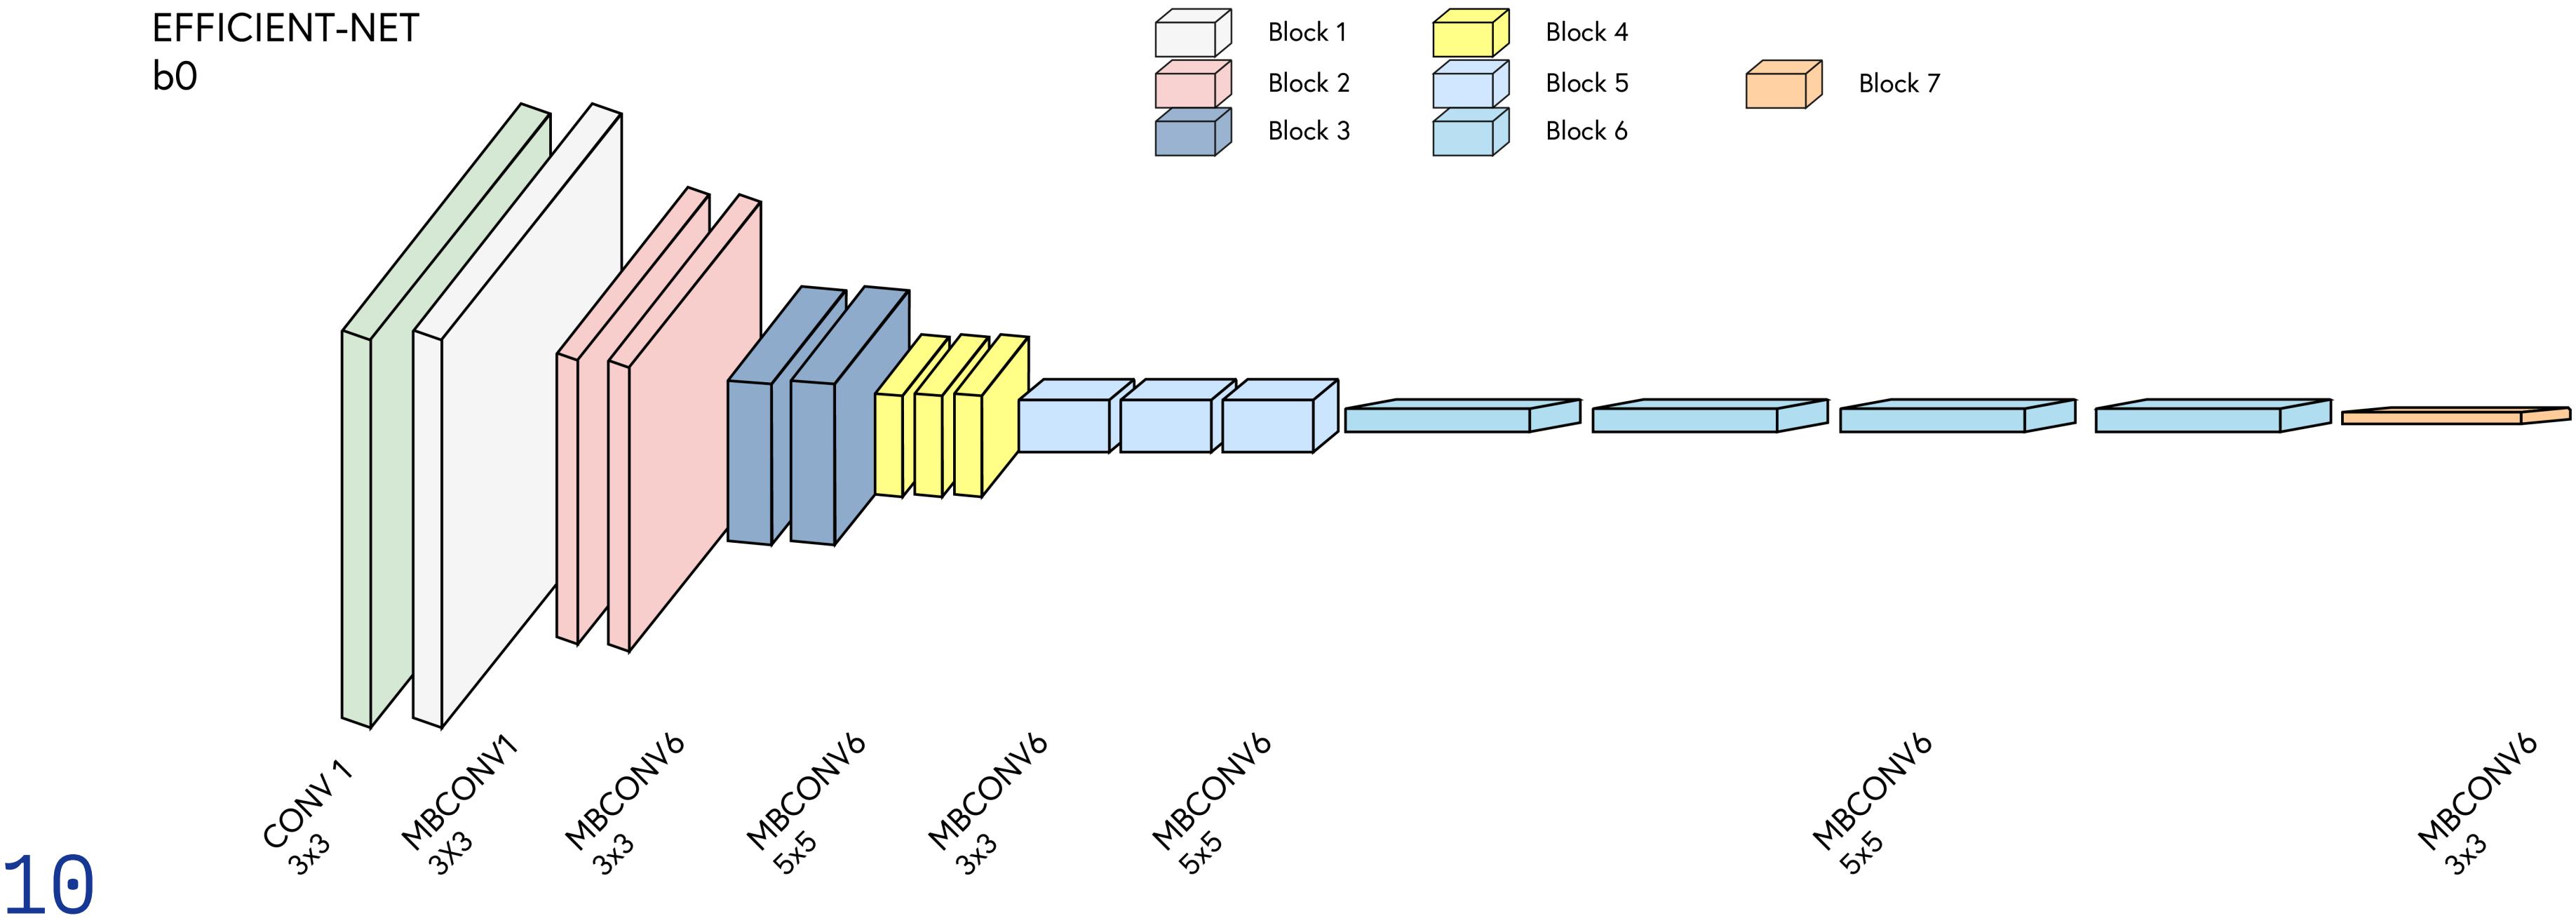
        # Calculate Intersection and Union
        intersection = (pred * target).sum()
        dice = (2. * intersection + self.smooth) / (pred.sum() + target.sum() + self.smooth)

    return 1 - dice # Return the loss to be minimized

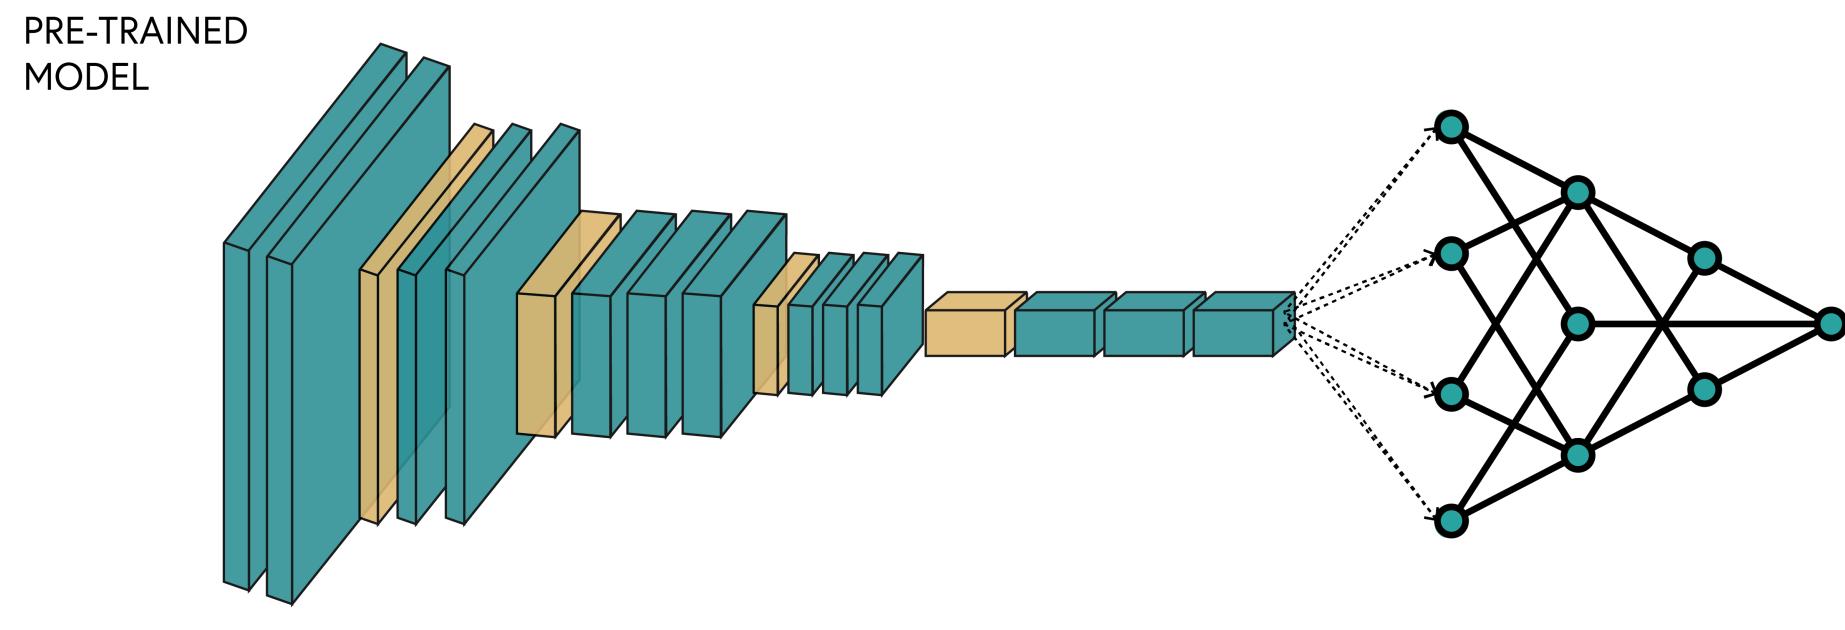
```

EFFICIENT-NET

introduced by google. achieves state-of-the-art accuracy while being significantly smaller and faster than previous models like ResNet or VGG. trained on millions of images (image-net), it already knows how to detect edges, textures, and complex shapes perfectly.



TRANSFER LEARNING



Step 1: Select a Pre-Trained Model

Choose a model architecture that has already been trained on a massive, diverse dataset

Step 2: Modify the Architecture

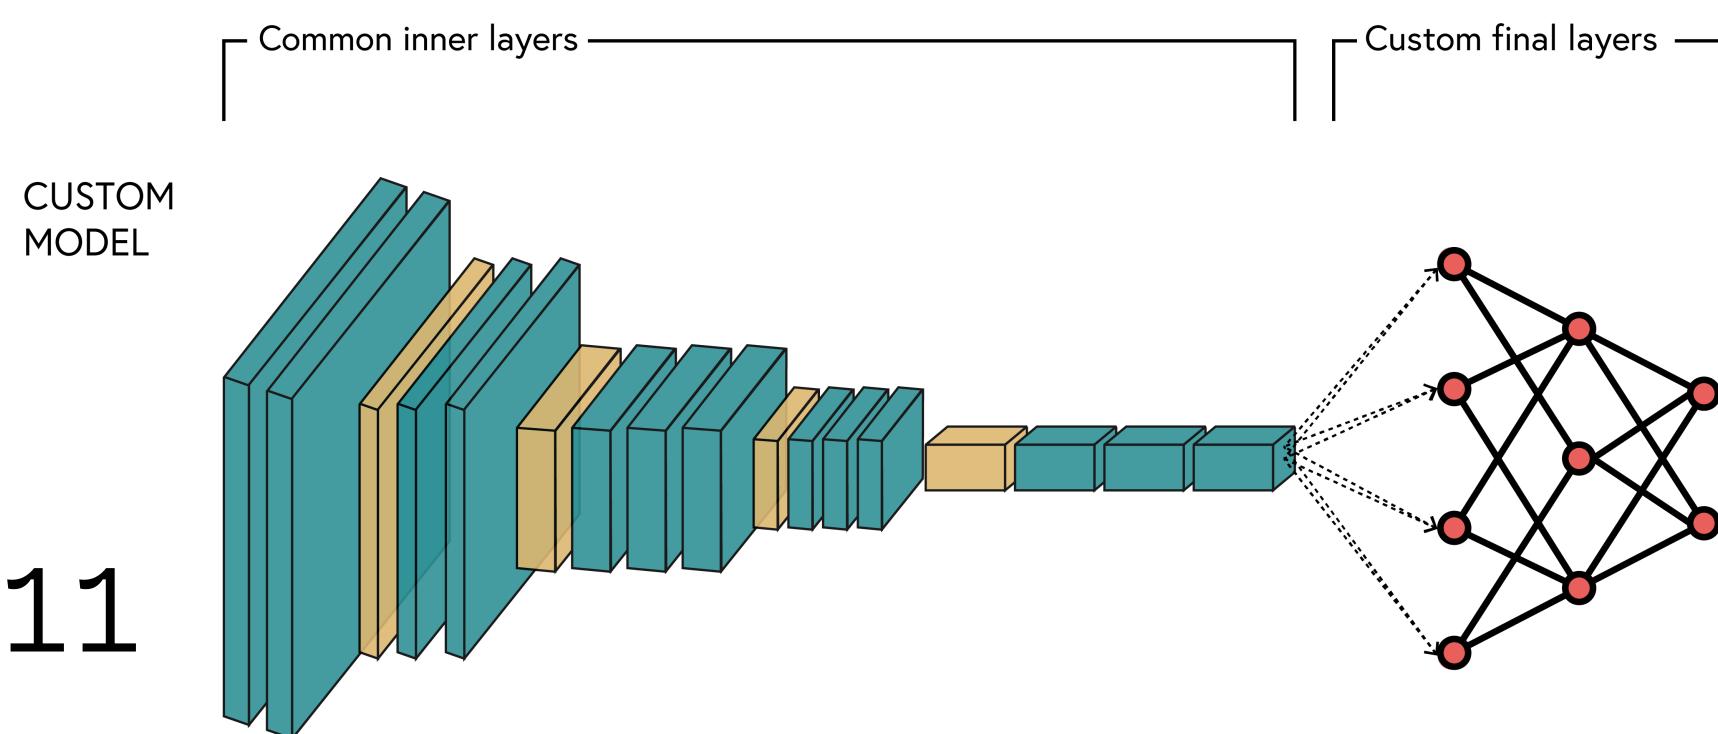
Strip away the original "head" of the model.

Step 3: Freeze the Base Layers

Lock the weights of the pre-trained layers (requires_grad = False).

Step 4: Add and Train New Layers

Attach your new, untrained layers tailored to your specific task (our U-Net Decoder).



Step 5: Fine-Tuning (Optional but powerful)

Once the new layers are stable, "unfreeze" the entire model. Train for a few more epochs using a very small learning rate.

PRE-TRAINED WEIGHTS

the `torchvision.models` API: pytorch makes it incredibly simple to download pre-trained models.

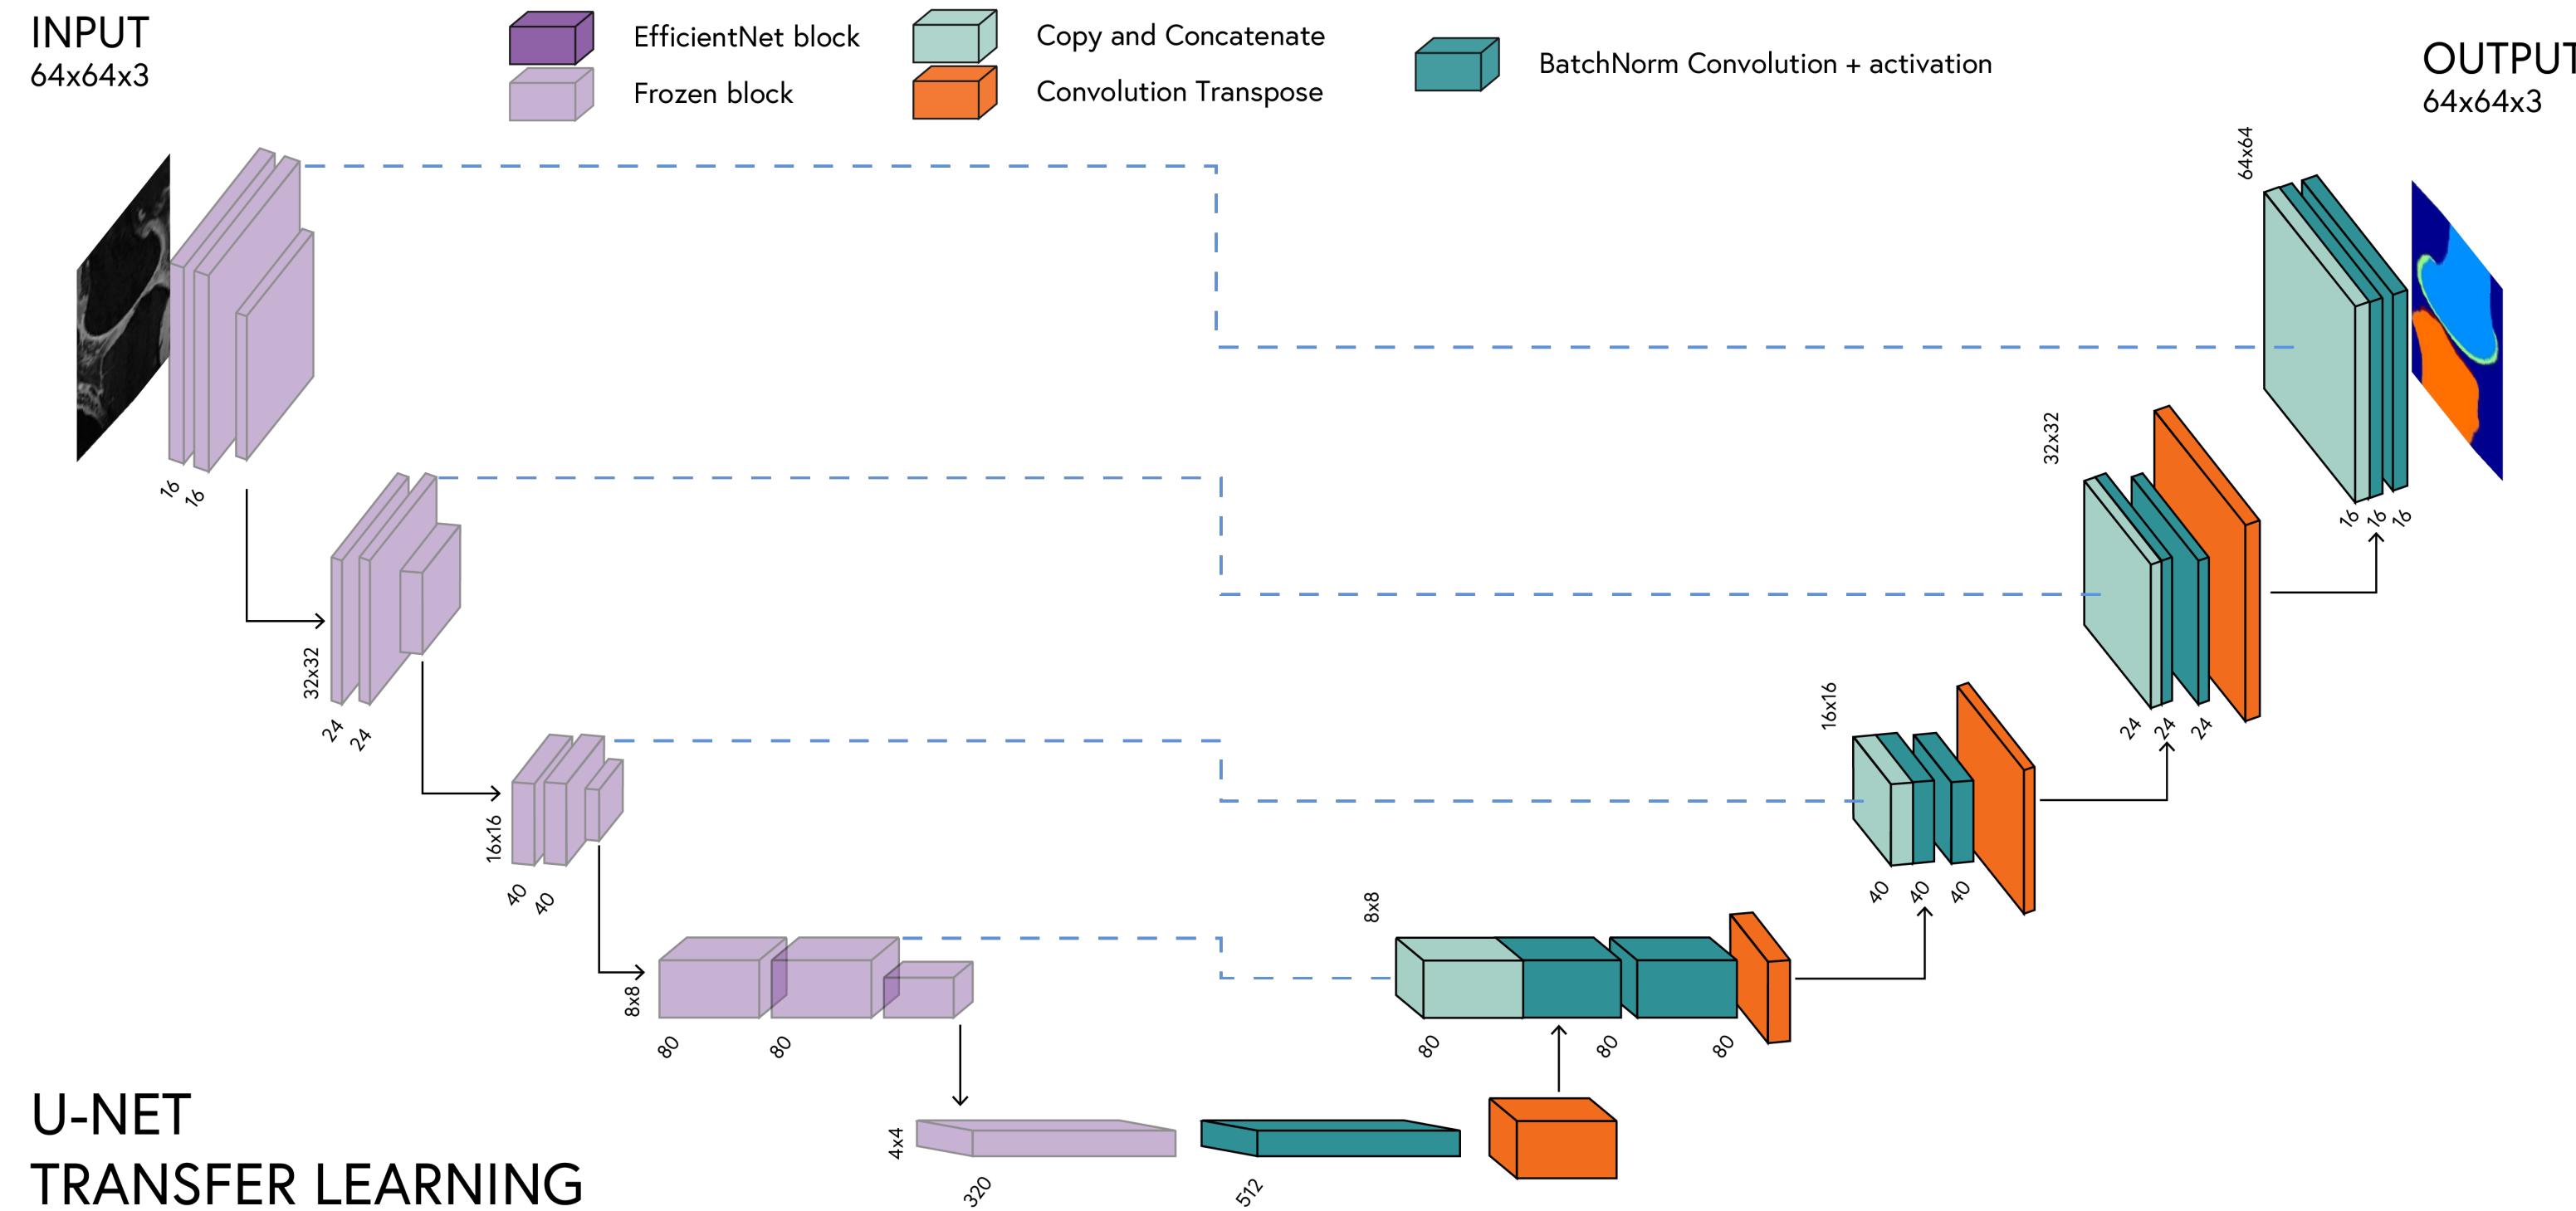
the `weights` object: we specify `IMAGENET1K_V1` to ensure we get the model that has already learned from the massive imagenet dataset.

feature extraction (`features`): we don't want the whole classification model. we specifically target the `.features` module of `efficientnet` to extract the intermediate feature maps needed for our u-net skip connections.

```
● ● ●  
import torchvision.models as models  
from torchvision.models import EfficientNet_B0_Weights  
  
# 1. Download the pre-trained EfficientNet-B0 model  
weights = EfficientNet_B0_Weights.IMAGENET1K_V1  
efficientnet = models.efficientnet_b0(weights=weights)  
  
# 2. Isolate the feature extractor (ignoring the classification head)  
encoder_features = efficientnet.features  
  
# Note: We will tap into specific blocks (e.g., layers 1, 3, 5)  
# to act as our skip connections for the U-Net Decoder!
```

EFFICIENT U-NET

the hybrid architecture: we are fusing two ideas: the powerful feature extraction of efficientnet (encoder) + the precise spatial localization of u-net (decoder).



FREEZING

the danger of random initialization:

our new u-net decoder has completely random weights. if we start training immediately, the massive errors from the random decoder will backpropagate into our pre-trained encoder.

the solution -> "freezing": we must turn off gradient tracking (`requires_grad = False`) for the efficientnet encoder during the initial training phase.

the workflow: train only the decoder for a few epochs until it learns how to use the encoder's features. then, optionally "unfreeze" the encoder and train everything together with a tiny learning rate (fine-tuning).



```
class EfficientUNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Load pre-trained encoder
        self.encoder = models.efficientnet_b0(
            weights=EfficientNet_B0_Weights.IMAGENET1K_V1).features

        # FREEZE the encoder weights!
        for param in self.encoder.parameters():
            param.requires_grad = False

        # Define the randomly initialized Decoder...
        # self.decoder_block1 = ...

    def forward(self, x):
        # ...
```

SMART-UP

standard u-nets are perfectly symmetrical.

however, when you swap in an external pre-trained encoder like efficientnet, the spatial dimensions (height/width) of the skip connections might not perfectly match your decoder's upsampled feature maps due to different padding or strides in the efficientnet blocks.

the error: if you try to concatenate two tensors with different spatial sizes (e.g., a 56x56 map with a 55x55 map), pytorch will throw a runtime shape error.

the solution: after the standard transposed convolution, we check if the sizes match. if they don't, we dynamically resize the upsampled tensor to perfectly match the skip connection using bilinear interpolation.

```
import torch
import torch.nn.functional as F

class SmartUpBlock(torch.nn.Module):
    # ... (__init__ defines self.up and self.conv) ...

    def forward(self, x1, x2):
        # x1: from previous decoder layer
        # x2: skip connection from encoder

        x1 = self.up(x1) # Initial upsampling

        # Check spatial dimensions (ignoring batch and channel dims)
        if x1.size()[2:] != x2.size()[2:]:
            # Force x1 to match the spatial size of x2
            x1 = F.interpolate(
                x1,
                size=x2.size()[2:],
                mode='bilinear',
                align_corners=False # Prevents edge artifacts
            )

        # Now it is safe to concatenate along the channel dimension!
        x = torch.cat([x2, x1], dim=1)

        return self.conv(x)
```